

Muhammad Khan

AI Engineer

✉ muhammadkhano8561@gmail.com ☎ 03019487954  linkedin.com/in/muhammad-khan-53a192357
 github.com/Muhammad-Khan-Khichi  muhammadkhan.my

SUMMARY

AI Engineer with practical experience across the **LLM** stack **LangChain**, **RAG**, vector databases, and multi-agent systems built on a foundation of classical ML, **Python**, and **FastAPI**. Shipped projects include a voice ordering agent platform (**LiveKit**) and a full stack AI job search SaaS.

EDUCATION

BS Computer Science

University of the Punjab

2024 – 2028

Lahore

PROFESSIONAL EXPERIENCE

AI Engineer Intern

Senarios Pvt Ltd 

06/2026 – 07/2026

Lahore, Pakistan

- Designed and deployed end-to-end RAG pipelines embedding generation, vector indexing (ChromaDB → Qdrant Cloud migration), and semantic retrieval to power context-aware LLM responses
- Integrated multiple LLMs (Mistral AI, GPT-4.1) via LangChain into production applications, including prompt-driven resume parsing and job matching, and real-time multi-agent voice AI systems combining speech-to-text, LLM reasoning, and text-to-speech

PROJECTS

HireAI (AI-Powered Job Search Platform) 

FastAPI, React, PostgreSQL (Neon), Qdrant (Cloud), LangChain, Mistral AI

- Built and deployed a **full-stack AI recruitment platform** with JWT + OAuth2 authentication, LLM-powered resume parsing, and AI-driven job matching, containerized via Docker on Hugging Face Spaces with a React frontend on Vercel
- Engineered a **resume-driven job search pipeline** using Mistral LLM for skill extraction, intelligent keyword generation, job match scoring/ranking, and content-hash-based deduplication to eliminate redundant LLM processing
- Developed **AI-powered career tools** (resume optimization, LinkedIn profile optimization, mock interviews, study plans, cover letters) and a **semantic search layer** with vector embeddings plus an admin dashboard for user, resume, and application analytics

AI Voice Agent for Restaurant Management

Python, LiveKit Agents, FastAPI, PostgreSQL, RAG, Deepgram Nova-3, Cartesia Sonic-3, Silero VAD

- Designed a **multi-agent AI voice assistant** using LiveKit with 5 specialized agents (Greeter, Reservation, Takeaway, Checkout, FAQ) for seamless restaurant operations with intelligent agent handoff logic
- Implemented **RAG-powered knowledge retrieval** indexing menu, FAQ, policies, and allergen data, and built a **real-time voice pipeline** integrating Silero VAD, Deepgram Nova-3 STT, OpenAI GPT-4.1 LLM, and Cartesia TTS with multi-language support
- Developed a **FastAPI backend** with PostgreSQL database for managing orders, reservations, feedback, and chat history with full CRUD REST APIs and interactive Swagger documentation

SKILLS

ML & Data

ML Algorithms, NumPy, Pandas, Matplotlib, Seaborn, Data Preprocessing, Feature Engineering

AI & LLM

LangChain, RAG, Livekit, Prompt Engineering, Vector Embedding

Backend & Deployment

Python, Fastapi, SQL, Qdrant DB, HuggingFace Spaces, Git/GitHub, Docker(basic)